

```

import os
import numpy as np
import random
import matplotlib.pyplot as plt
import image_resources as ir
from decimal import *

# Functions that might be useful (please read the documentation)
# x.flatten() (take a N-dimensional numpy array and make it one-
dimensional)
# numpy.random.choice -- choose from the list of images
# numpy.dot -- compute the dot product
# numpy.random.normal -- set up random initial weights

DIM = (28,28) #these are the dimensions of the image

def load_image_files(n, path="images/"):
    # helper file to help load the images
    # returns a list of numpy vectors
    images = []
    for f in os.listdir(os.path.join(path,str(n))): # read files in
the path
        p = os.path.join(path,str(n),f)
        if os.path.isfile(p):
            i = np.loadtxt(p)
            assert i.shape == DIM # just check the dimensions here
            # i is loaded as a matrix, but we are going to flatten it
into a single vector
            images.append(i.flatten())
    return images

## Your code here:
def saveImages(files, show=False):
    """Saves digit image matrices for faster access during training"""

    for j in range(len(files)):
        np.save(files[j], load_image_files(j))
        if show:
            print(files[j] + '.npy saved')

def loadImages(files, show=False, unseen=False):
    """Returns a dictionary of image files with digits as keys"""
    trainingDict = dict()
    unseenDict = dict()

    for f in range(len(files)):
        trainingDict[f] = np.load(files[f] + '.npy')
        if show:
            print(files[f] + '.npy loaded')

```

```

        if unseen:
            print('Creating a dictionary of test images, 500 for each
digit...')
            for x in range(len(trainingDict)):
                unseenDict[x] = np.empty((500, 784))

                for i in range(500):
                    randomIndex = random.randint(0, len(trainingDict[x]) -
1)
                    unseenDict[x][i] = trainingDict[x][randomIndex]
                    trainingDict[x] = np.delete(trainingDict[x],
randomIndex, 0)
            print('...')
            print('Test dictionary complete')

```

```

    return trainingDict, unseenDict

```

```

fileNames = ['Zero', 'One', 'Two', 'Three', 'Four', 'Five', 'Six',
'Seven', 'Eight', 'Nine']

```

```

trainingImages, unseenImages = loadImages(fileNames, unseen=True)
N = len(trainingImages[0][0])
assert N == DIM[0]*DIM[1] # just check our sizes to be sure

```

"""Problem 1: Write an implementation of the perceptron learning
 algorithm that first
 loads images for the digit "0" and then for the digit "1". Start
 with random weights
 from a normal distribution. Compute the average accuracy on blocks
 of 25 items and plot
 this accuracy until you think it won't get better.

ANSWER: The answer to how many images it takes to train my model
 varies depending
 on the random weights I start with but it usually takes around
 400 blocks to
 achieve achieve .99 accuracy. This is about 10,000 images
 between the 0 and 1
 training sets. I attempted to get .9999 accuracy but it takes
 10x longer
 and doesn't seem to produce significantly better results. """

```

class Perceptron(object):

```

```

    def __init__(self, dimensions, data, digit0, digit1):
        self.weights = np.random.normal(0, 1, size=dimensions)
        self.data_set = {0: data[digit0], 1: data[digit1]}
        self.digits = [digit0, digit1]
        self.bias = np.zeros(dimensions)

```

```

        self.overall_accuracy = 0

    def get_weights(self):
        return self.weights

    def dotProd(self, w, x):
        return np.dot(w, x)

    def predict(self, w, x, b):
        dot = self.dotProd(w, x)

        if dot >= 0:
            return 1

        else: return 0

    def classify(self, w, x):
        dot = self.dotProd(w, x)

        if dot >= 0:
            return 1
        else: return 0

    def random_label(self):
        """Returns a random dataset label"""
        labels = list(self.data_set.keys())
        return random.choice(labels)

    def plot_accuracy(self, x, y, save=False, name=None, color='red'):

        plt.xlabel('Blocks of 25')
        plt.ylabel('Accuracy')
        plt.title('P1: Training Accuracy for Digits {0} and
{1}'.format(self.digits[0], self.digits[1]))
        plt.legend(loc=0, title='Converged to {0} Accuracy in {1}
Blocks'.format(round(y[-1], 5), x[-1]))
        plt.plot(x, y, c=color, label='')

        if save:
            plt.savefig(name + '.pdf')

        plt.show()
        plt.close()

    def train(self, threshold, precision, blocks, plot=False):
        """Data_set is a dictionary of lists containing (784,)
arrays"""
        accuracy = .5
        all_blocks = blocks

```

```

correct = 0
accuracy_trace = []
xAxis = []
acc_delta = 1
getcontext().prec = precision

while (all_blocks // 25) <= 500:

    xAxis += [all_blocks // 25]
    acc_delta = accuracy

    for i in range(blocks):
        label = self.random_label()
        data = self.data_set[label]
        sub_data = random.choice(data)
        y = self.predict(self.weights, sub_data, self.bias)

        if label == 0 and y == 1:
            self.bias -= sub_data
            self.weights -= sub_data - self.bias
        elif label == 1 and y == 0:
            self.bias += sub_data
            self.weights += sub_data + self.bias
        else:
            correct += 1

    accuracy = correct / all_blocks
    acc_delta = abs(round(Decimal(accuracy), precision) -
round(Decimal(acc_delta), precision))
    all_blocks += blocks
    accuracy_trace.append(accuracy)

self.overall_accuracy = accuracy

if plot:
    self.plot_accuracy(xAxis, accuracy_trace)

# zeroOnePercept = Perceptron(N, trainingImages, 0, 1)
# zeroOnePercept.train(.97, 5, 25, plot=True)
#
# np.save('zero_one_weights', zeroOnePercept.get_weights())

```

"""Problem 2: Does your solution in Q1 converge on 100% accuracy or not? What does this mean in terms of the linear separability of "0" and "1" on this feature space?

ANSWER: My solution for Q1 approaches 100% (I used the method of total correct / total images) which reaches .9999 for digits 0 and 1. This tells

me that 0 and 1  
are completely linearly separable. The pair have unique features  
that allow for  
distinguishing them completely. """

"""Problem 3: What do large negative and large positive values mean,  
intuitively? What  
do numbers near zero mean? Why does this matrix look the way that  
it does, in terms  
of where large positive and negative terms are located?

ANSWER: The large positive and negative values seem to occur where  
the digits were  
in the training images. For my perceptron, 1 took very large  
values while 0 took  
very negative values. In my matrix, the general shape of a 0 can  
be seen in black  
which is the color for large negative values and bright yellow can  
be seen in the center  
of the 0 where a 1 would be if the digits overlapped. There is  
also yellow outside the  
black 0 shape and I take this to represent some of the 1 images  
that were written in  
a 'forward slash' shape. Numbers near 0 represent the 0's (as  
opposed to the 1's) from  
the training images where the weights were updated but the feature  
values in that area  
was 0 so the weights stayed small."""

```
def weightMatrix(weights, dims, save=False,
fileName='FILENAME',method=None, bounds=[0, 1]):
    # weights= weights[10]
    weights = weights.reshape(dims, dims)
    fig, im = plt.subplots(figsize=(10, 10)) #
    wm = plt.imshow(weights) #, vmin=bounds[0], vmax=bounds[1]
    #cmap=plt.get_cmap(ir.color_maps[ir.color_maps.index('inferno')]),
    interpolation=ir.interpol_methods[method],
    wm.axes.get_xaxis().set_visible(False)
    wm.axes.get_yaxis().set_visible(False)
    plt.title('Heat Map of Trained Weight Matrix For Digits 0 and 1',
fontSize=15)
    fig.colorbar(wm, orientation='horizontal', fraction=.0415)
    if save:
        plt.savefig(fileName + '.pdf')

    plt.show()

# boundList = [0, 1]
# weightMatrix(np.load('zero_one_weights.npy'), 28, save=True,
fileName='a7p3_norm_bias', method=0, bounds=boundList)
```

```
"""Problem 4: What should you expect to happen if you set the elements
of the weight vector
    which are close to zero to be actually zero? Do this for the 10,
20, 30, ... 780 weight
    values closest to zero (in absolute value) and plot the resulting
accuracies on 1000 random
    classifications of "0" vs "1". What does this tell you about the
proportion of the image
    which is diagnostic about "0" vs "1"?
```

```
ANSWER: Setting the weights 'close' to 0 should negate the effect
of those elements; they
    shouldn't contribute to the classification of the digits. My plot
showed near 100% classification
    accuracy up to about 120 elements set to 0 and then the accuracy
falls to and hovers around 50%.
    When the accuracy falls to 50% the perceptron is essentially
guessing as it doesn't have enough
    feature data to know which digit is which.
```

```
Seeing that about 121 elements of the weight array can essentially
be ignored tells me that
    a low proportion of the image is diagnostic about 0 vs 1
especially when you consider
    that a large portion of the weights were not effected by training
in the first place.
    I created a matrix of the weight array using the last version with
~98% accuracy and
    it contained only 4 or 5 elements in the area of the digits but it
was enough for linear
    separation.
"""
```

```
def testClassification(numIters, digit0, digit1, perceptron,
unseenDict, weights):
    numCorrect = 0

    for i in range(numIters):
        label = random.choice([digit0, digit1])
        result = perceptron.classify(weights, unseenDict[label][i %
(numIters // 2)])

        if label == digit0 and result == 0:
            numCorrect += 1
            label = digit1
        elif label == digit1 and result == 1:
            numCorrect += 1
            label = digit0
```

```

        return numCorrect

def setToZero(trainDict, unseenDict, N, digit0, digit1):
    accuracyList = []
    weightsList = []
    xList = [10 * i for i in range(1, 79)]
    perceptron01 = Perceptron(N, trainDict, digit0, digit1)
    perceptron01.train(.97, 5, 25)
    modWeights = np.copy(perceptron01.get_weights())
    modWeights = np.absolute(modWeights) #convert all to positive
values
    modWeights[modWeights == 0] = 1000 #make all 0's into 1000, if any

    for x in range(1, 79):
        k = 0

        while k < (len(range(x * 10))):
            index = np.argmin(modWeights)
            perceptron01.weights[index] = 0
            modWeights[index] = 1000
            k += 1

        result = testClassification(1000, digit0, digit1,
perceptron01, unseenDict, perceptron01.weights)
        accuracyList += [result / 1000]
        weightsList += [np.copy(perceptron01.weights.reshape(28, 28))]

    return accuracyList, xList, weightsList

# accList, xList, moddedWeights = setToZero(trainingImages,
unseenImages, N, 0, 1)
# plt.plot(xList, accList)
# plt.title('Modified Weight Values at Intervals 10, 20, 30, ...,
780')
# plt.ylabel('Average Accuracy')
# plt.xlabel('Weight Values Changed to 0 (smallest first)')
# plt.savefig('a7p4.pdf')
# plt.show()
# plt.close()
#
# plt.imshow(moddedWeights[10])
# plt.savefig('mod_at_11.pdf')
# plt.close()
# plt.imshow(moddedWeights[13])
# plt.savefig('mod_at_13.pdf')
# plt.close()

```

"""Problem 5: Next show a matrix of the classification accuracy of each pair of digits

after enough training. Make this a plot (with colors for accuracy rather than numbers).

Does it match your intuitions about which pairs should be easy vs. hard? Why or why not?

ANSWER: The matrix does approximately match my intuition. All digits compared with themselves (0 and 0 for example) have 100% accuracy. The perceptron is 96 – 100% accurate for most pairs. I would have expected (4, 9), (3, 5), and (5, 8) to be harder to linearly separate due to the amount of overlap these digits tend to have in handwriting and my matrix reflects this intuition with 90–92% accuracy for these digit pairs.

"""

```
def allDigitsAcc(trainDict, unseenDict, N):
    digitArray = np.zeros(100)
    digitArray = digitArray.reshape(10, 10)
    xlabels = np.arange(0, 10)
    ylabels = np.arange(0, 10)

    for k in range(len(trainDict.items())):
        for j in range(len(trainDict.items())):
            correct = 0
            percept = Perceptron(N, trainDict, k, j)
            percept.train(.97, 5, 25)
            result = testClassification(1000, k, j, percept,
unseenDict, percept.get_weights())
            digitArray[k, j] = result / 1000

    return digitArray, xlabels, ylabels

#uncomment below to run code for Problem 5
# grid, xAx, yAx = allDigitsAcc(trainingImages, unseenImages, N)
# np.save('accMatrix', grid)
# np.save('accMatrix_X', xAx)
# np.save('accMatrix_Y', yAx)
#
grid = np.load('accMatrix.npy')
xAx = np.load('accMatrix_X.npy')
yAx = np.load('accMatrix_Y.npy')

plt.rcParams['xtick.bottom'] = plt.rcParams['xtick.labelbottom'] =
False
plt.rcParams['xtick.top'] = plt.rcParams['xtick.labeltop'] = True

fig, ax = plt.subplots() #figsize=(10,10)
matrix = plt.imshow(grid, cmap='inferno', extent=[0, 9, 9, 0])
```



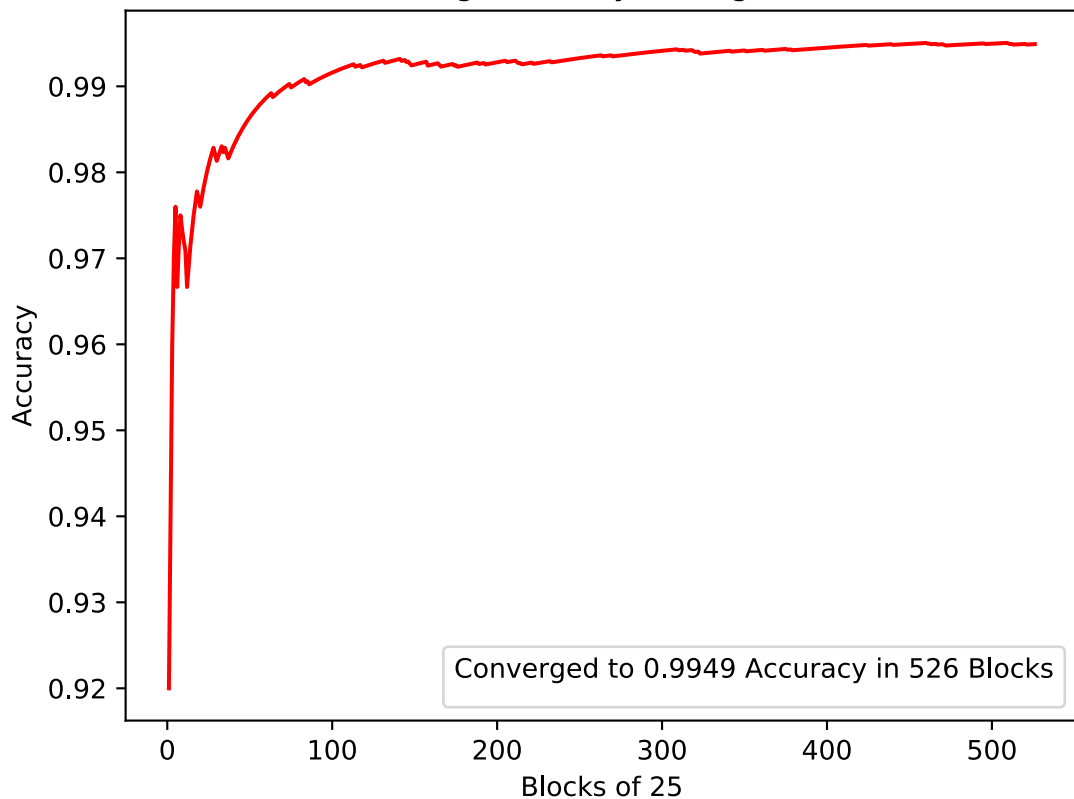
```
#'coolwarm'
figure_title = 'Classification Accuracy For Each Pair of Digits'
plt.text(0.5, 1.13, figure_title,
        horizontalalignment='center',
        fontsize=12,
        transform = ax.transAxes)
fig.colorbar(matrix, orientation='horizontal', fraction=.05)
ax.set_xticks(xAx)
ax.set_yticks(yAx)

plt.savefig('a7p5accuracy_4.pdf')


plt.show()

#end
```

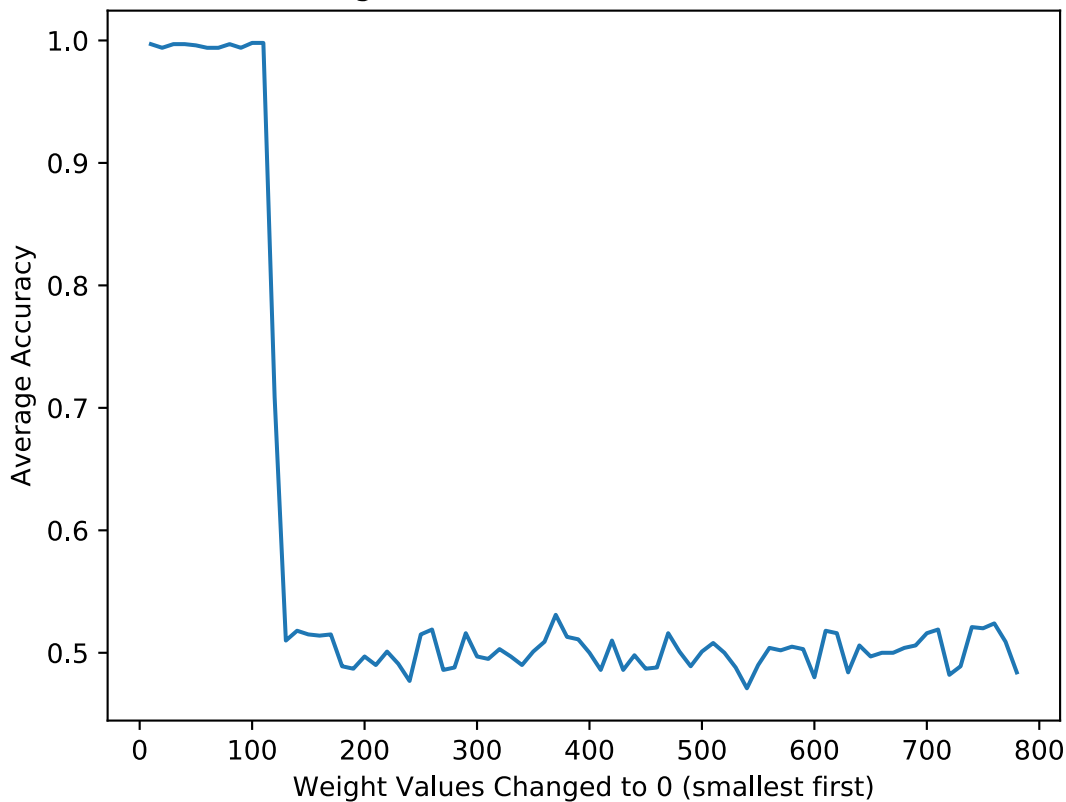
P1: Training Accuracy for Digits 0 and 1



Heat Map of Trained Weight Matrix For Digits 0 and 1



Modified Weight Values at Intervals 10, 20, 30, ..., 780



Classification Accuracy For Each Pair of Digits

